Agenda

- Course Overview
- Course Objectives and Outcomes
- Quick Review of Complete Handout
- Practical Aspects
- Different Types of Data and Storage for Data
- Big Data Characteristics, sources
- Big Data Systems Perspective – in-memory vs storage vs network
- Big Data challenges, applications/case studies
- Locality of Reference – principle examples

BIG DATA ANALYTICS – emphasis / important ones

Introduction to Industry 4.0/5.0, BI & Analytics
Big Data Definition & Characteristics
Sources of Big Data
Challenges & Benefits of Big Data Systems
Case Studies / Applications

Hadoop Ecosystem & Map Reduce
Hadoop Intro
HDFS architecture
MAP REDUCE architecture
SPARK Usage

HIVE Database & Others
Hive Architecture and uses
Pig, Sqoop, Flume, OOzie
Other popular tools (open source & commercial)

Mongodb Architecture & Commands
Data Visualization;
Applications/projects – PowerBI/Tableau & Excel

PRACTICAL / LAB EXAMPLES USING LINUX, SHELL, DB, SQL, JAVA, SQL, PYTHON, POWERBI.
MCQ – PUZZLES – CASE STUDIES/DISCUSSIONS

Focus Areas/Practicals

✔*VIRTUAL MACHINE*
✔*UBUNTU LINUX COMMANDS*
✔*EDITING FILES AND ENVIRONMENT*
✔*SETTING UP HADOOP CLUSTER*
✔*STARTING CLUSTER*
✔*PRACTICING HADOOP AND HDFS COMMANDS*
✔*SETTING UP AND RUNNING MAP REDUCE*
✔*SETTING UP HIVE DB*
✔*RUNNING QUERIES AGAINST HIVE DB*
✔*EXPORT/IMPORT TO HIVE*
✔*WORK WITH MONGODB*
✔*QUICK DEMO… OF OTHER TOOLS…*

A QUICK REVIEW OF DATA ANALYST & SCIENTIST

Data Analysts and Scientists
Job Profile: Data analysts help businesses develop well-informed strategies by creating charts and prepare visual presentations. Also, examining large data and identifying trends are the expected roles of a data analyst. Data scientists construct and develop new processes for data modelling and primarily use prototypes, algorithms, predictive models and custom analysis.
Chief skills: Python coding, Hadoop Platform, R programming, SQL Database/Coding, Apache spark, Machine learning and AI, Data Visualisation
Big Data Specialists
Job Profile: Utilise data analysis to evaluate the technical performance of an organisation. Also provides recommendations on system enhancements.
Job skills: Apache Hadoop, Apache Spark, NoSQL, Machine learning and Data Mining, Statistical and Quantitative Analysis, SQL, Data Visualization, General Purpose Programming language
Digital Transformation Specialists
Job Profile: Work in enhancing a company's technical performance. They analyse the company's infrastructure and the gaps in service.
Job skills: Technical aptitude, critical thinking abilities, excellent communication skills, adaptability, SQL, C++, HTML, CSS

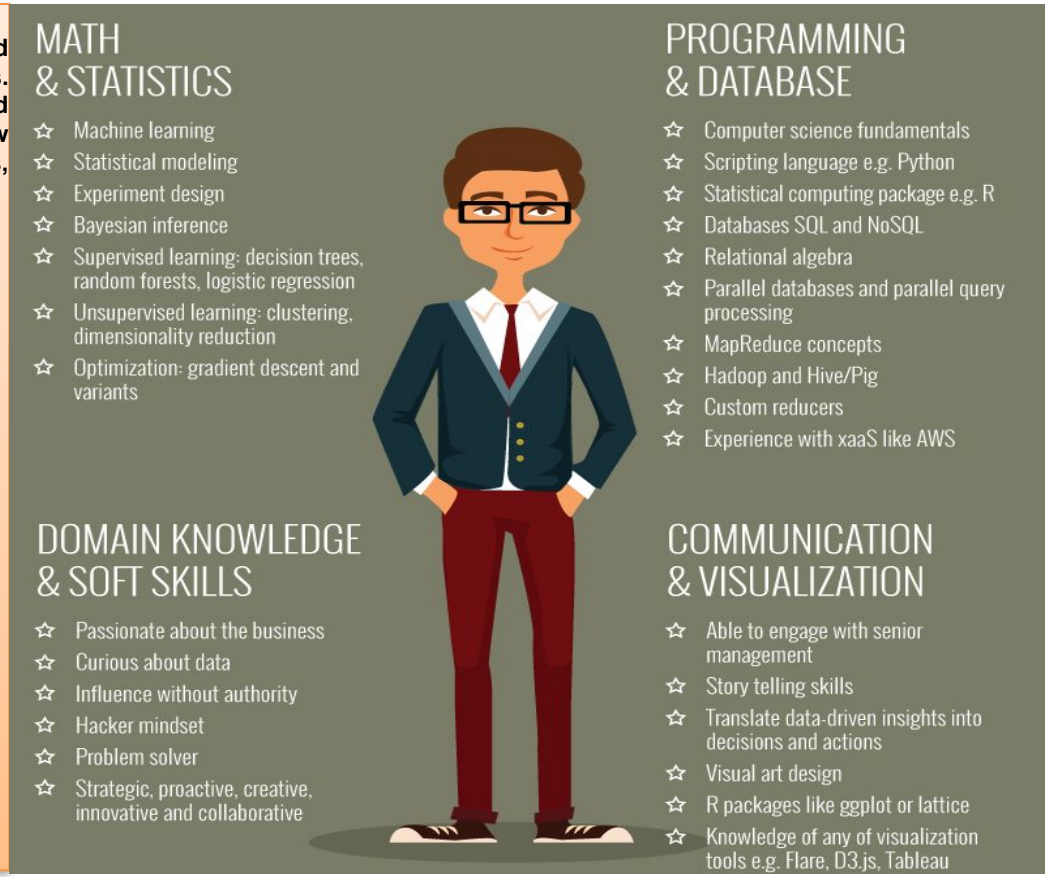

A FEW million $ QUESTIONS

- Data is woven into the everyday fabric of our lives. With the rise of mobile, social media, and smart technologies associated with the Internet of Things (IoT), we now transmit more data than ever before—and at a dizzying speed. Thanks to big data analytics!!
- Organizations can now use that information to rapidly improve the way they work, think, and provide value to their customers.
- With the assistance of tools and applications, big data can help you gain insights, optimize operations, and predict future outcomes.

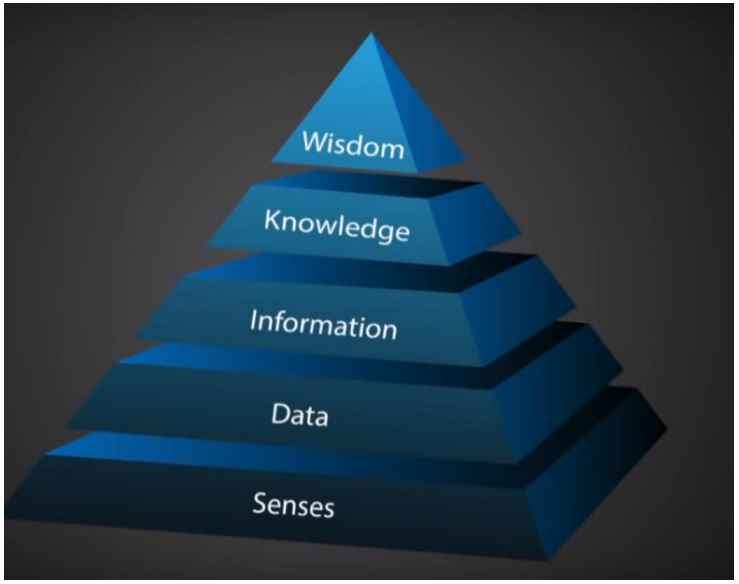

Information	Internet Of Everything/Things
Science	Basics/Formulas/neurons..
Engineering	Structural Approaches/principles/methods
Technology	Enablers… Wifi.. 5S.. etc
Intelligence	AI, ML, DL, ANN
Automation	Robots, Chatbots, Programs, Daemons, etc
Machine Intelligence	Making machine to think.. and communicate with other machines as well as human
B-B C-B..!? Management	Data Analytics, Visualization, Statistics

Societal Values & Cost Savings & New Way of Lazy Life??

EVOLUTION OF INDUSTRIES

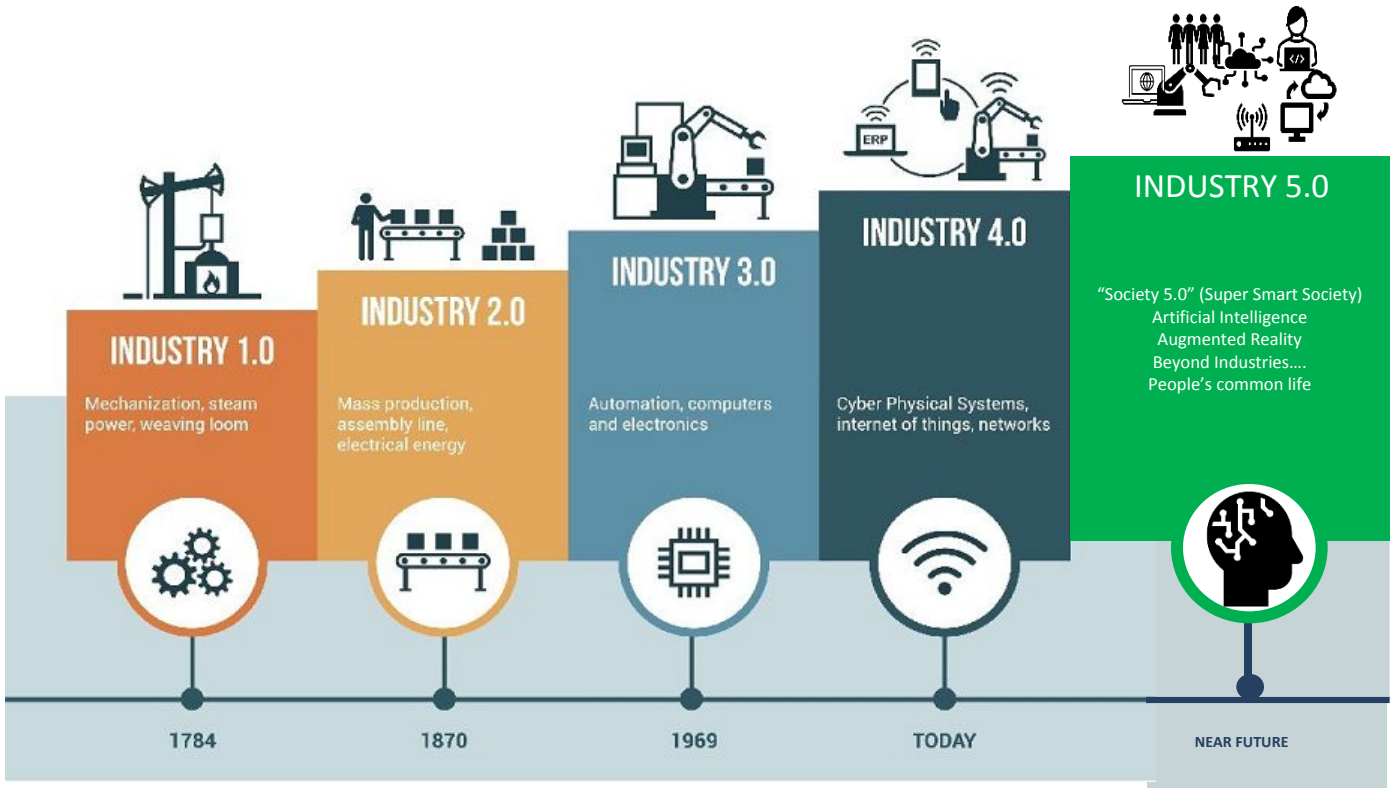

BUILDING BLOCKS OF INDUSTRY 4.0 & 5.0

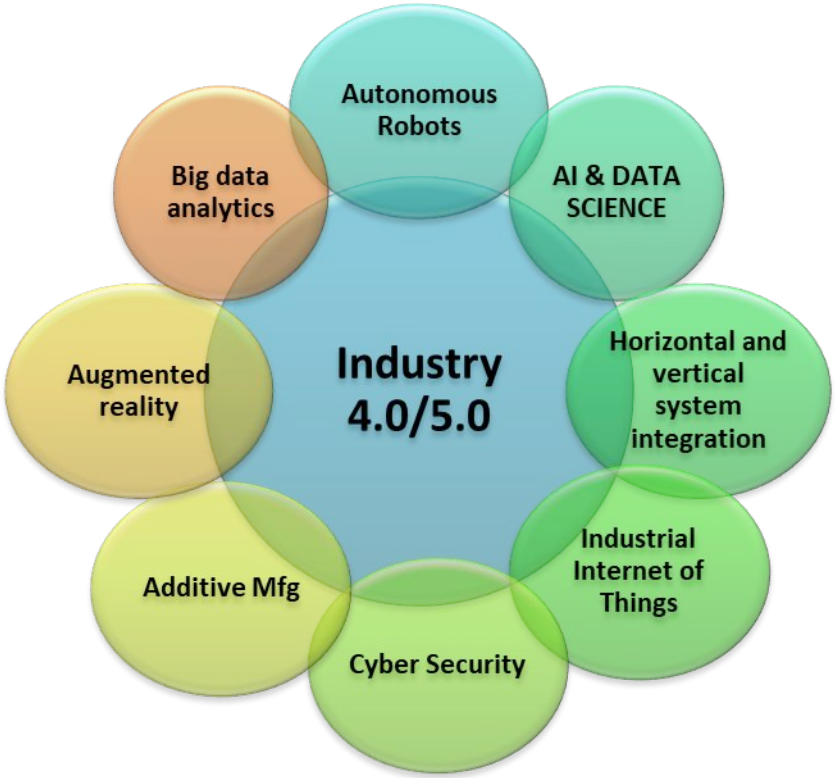

All the building blocks create a huge opportunities to the universities/engineering institutes to prepare or update the curriculum and teaching methodologies to train the engineers to be industry 4.0 read & 5.0 ready.

What Is Business Intelligence & How it works?

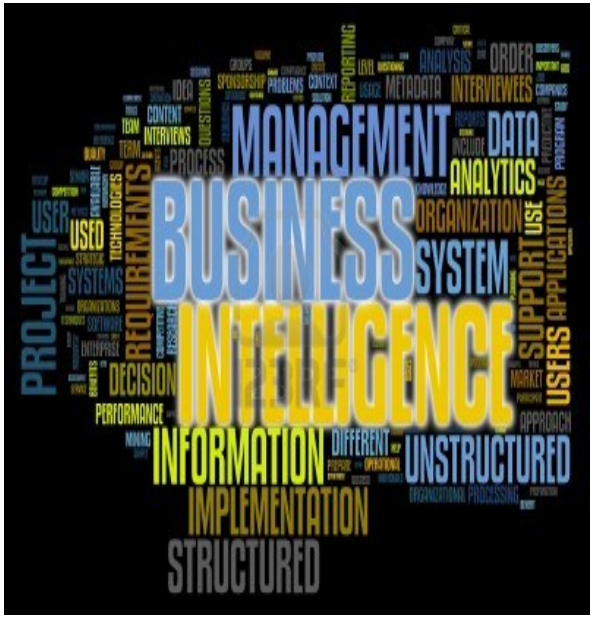

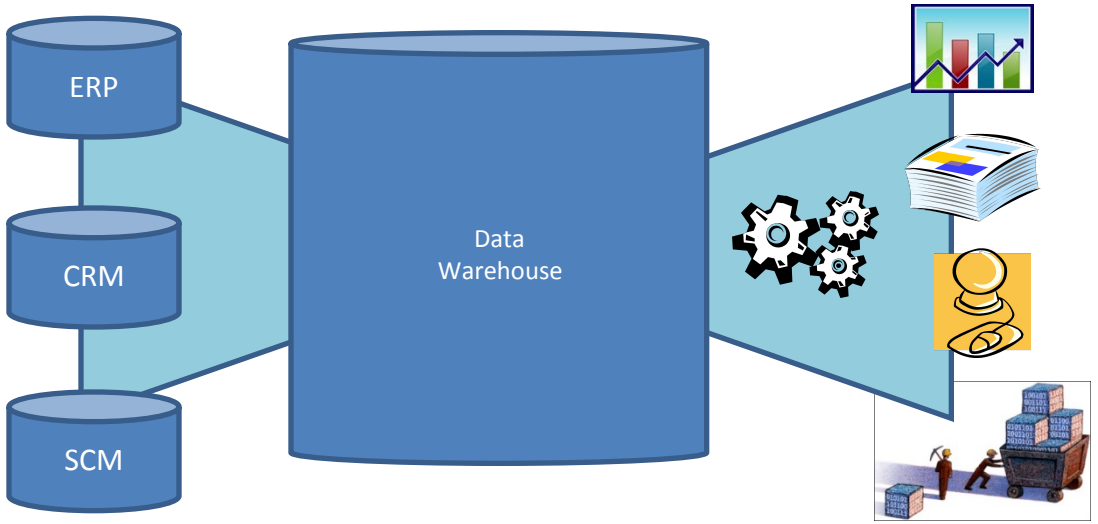

IBM COGNOS

SAP BI

Microsoft Business Intelligence

"Getting data in"

"Getting data out"

The Scope of Business Intelligence

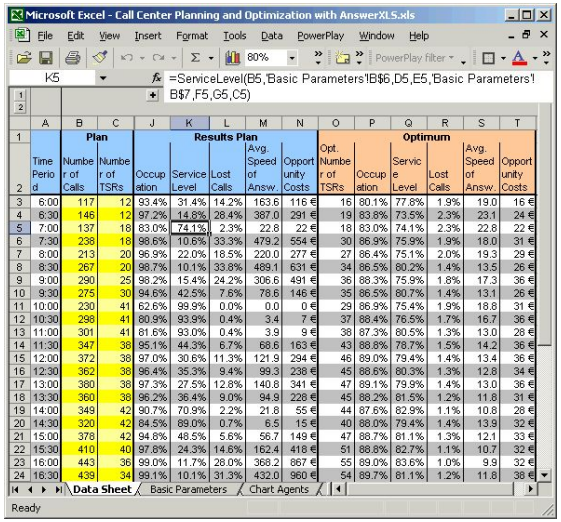

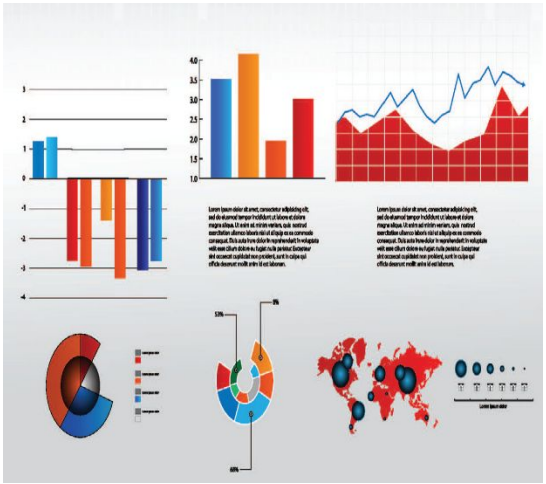

Smaller organizations:
Excel spreadsheets

Larger organizations:
Data mining, Predictive, Prescriptive, analytics, dashboards

WHO IS DATA SCIENTIST ??

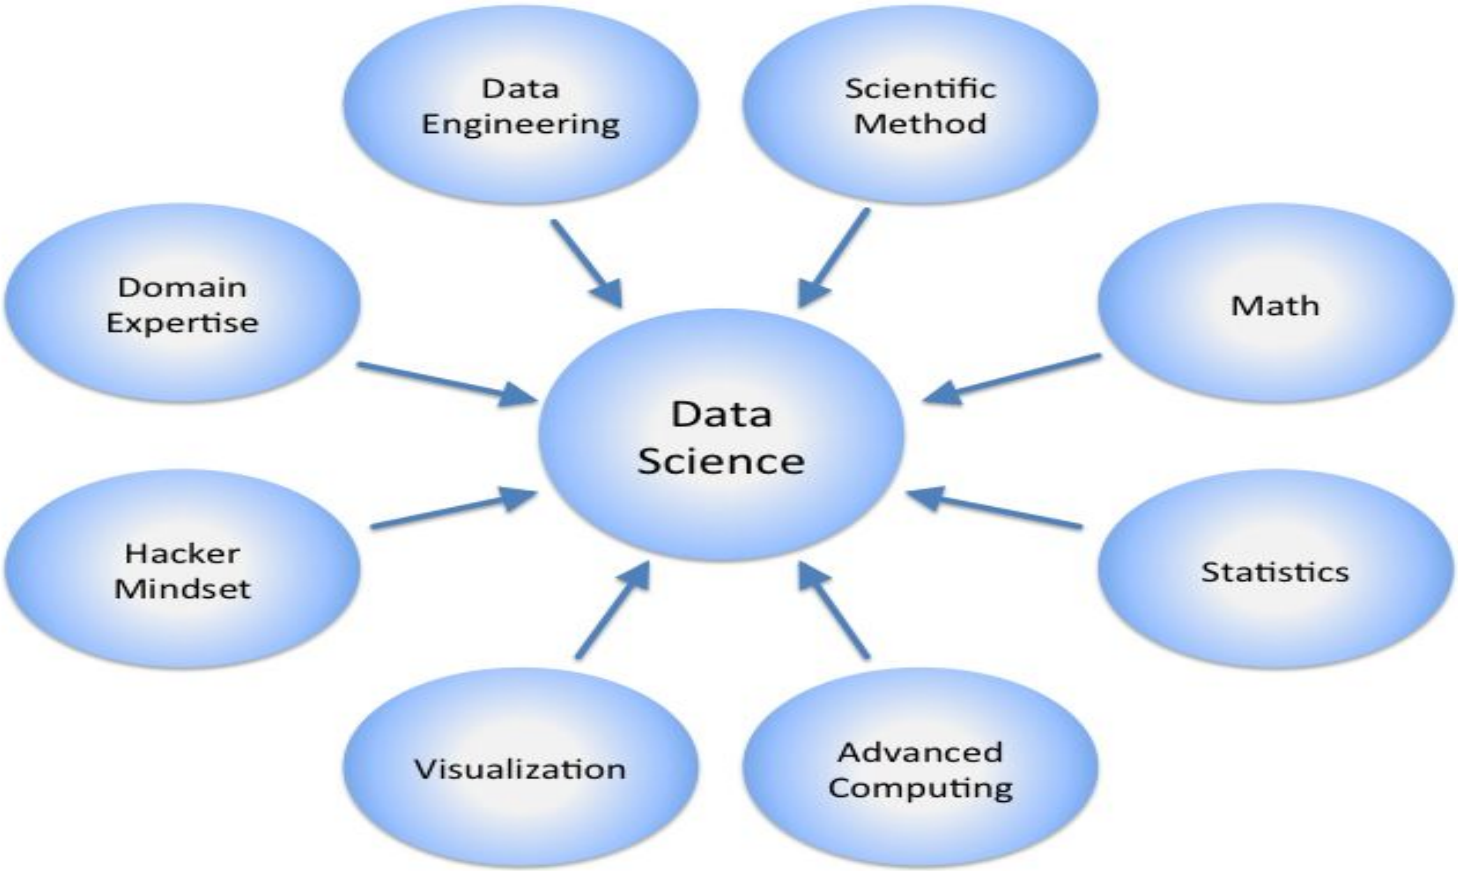

- There are two basic types of Predictive Analytics / Data Science problems:
- 1. Internal Predictive Analytics / Data Science problems, such as bad data, reckless analytics, or using inappropriate techniques.
- Internal problems are not business problems; they are internal to the Predictive Analytics / Data Science community.
- Therefore, the fix consists in training data scientists to do better work and follow best practices.
- 2. Applied business problems are real-world problems for which solutions are sought, such as fraud detection or identifying if a factor is a cause or a consequence.
- These may involve internal or external (third-party) data.

- These are the characteristics of the modern trends in Predictive Analytics / Data Science which one you should be aware of:

  - **In-memory analytics**
  - MapReduce and Hadoop
  - NoSQL, NewSQL, and graph databases
  - Python and R
  - Data integration: blending unstructured and structured data (such as data storage and security, privacy issues when collecting data, and data compliance)
  - Visualization
  - Analytics as a Service, abbreviated as AaaS
  - Text categorization/tagging and taxonomies to facilitate extraction of insights from raw text and to put some structure on unstructured data

- What Knowledge does a Data Scientist need:

- Thus, data scientists also need to be good communicators to understand, and many times guess, what problems their client, boss, or executive management is trying to solve.

- Translating high-level English into simple, efficient, scalable, replicable, robust, flexible, platform-independent solutions is critical.

- Predictive Analytics / Data Science = Some (computer science) + Some (statistical science)
- + Some (business management) + Some (software engineering) + Domain
- expertise + New (statistical science), where

- Some () means the entire field is not part of Predictive Analytics / Data Science.
- New () means new stuff from the field in question is needed

- Horizontal Versus Vertical Data Scientist
- Vertical data scientists have deep technical knowledge in some narrow field.
- For instance, they might be any of the following:
- Computer scientists familiar with computational complexity of all sorting algorithms
- Statisticians who know everything about eigenvalues, singular value decomposition and its numerical stability, and asymptotic convergence of maximum pseudo-likelihood estimators
- Software engineers with years of experience writing Python code (including graphic libraries) applied to API development and web crawling technology

- Horizontal Versus Vertical Data Scientist

- Database specialists with strong data modeling, data warehousing, graph database, Hadoop, and NoSQL expertise

- Predictive modelers with expertise in Bayesian networks, SAS, and SVM

- The key here is that by “vertical data scientist” we mean those with a more narrow range of technical skills,
- such as expertise in all sorts of Lasso-related regressions but with limited knowledge of time series, much less of any computer science.

- Horizontal Versus Vertical Data Scientist

- Horizontal data scientists are a blend of business analysts, statisticians, computer scientists, and domain experts.

- They combine vision with technical knowledge.

- They might not be experts in eigenvalues, generalized linear models, and
- other semi-obsolete statistical techniques,
- but they know about more modern data-driven techniques applicable to unstructured, streaming, and big data.

- They can design robust, efficient, simple, replicable, and scalable code and algorithms.

Data Science

- Statistical and Operations research techniques
- Machine Learning
- Deep Learning

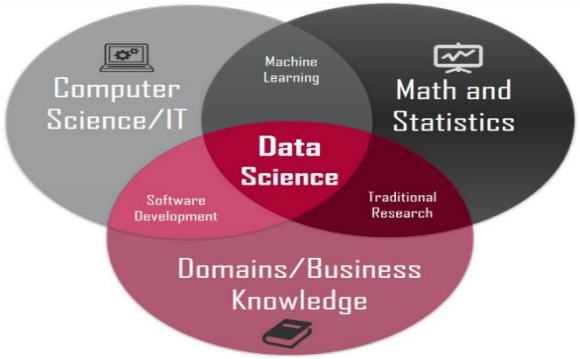

Data Science refers to an emerging area of work concerned with the **Collection, Preparation, Analysis, Visualization, Management, and Preservation** of **large collections of information**.

Data Science involves using methods to analyze **massive amounts of data** and **extract the knowledge** it contains.

Data Science includes **Data Analysis** as an important component of the skill set required for many jobs in this area, but is not the only necessary skill.

Data Science is fundamentally an interdisciplinary subject

Data Science involves using methods to analyze **massive amounts of data (Big Data)** and extract the knowledge it contains.

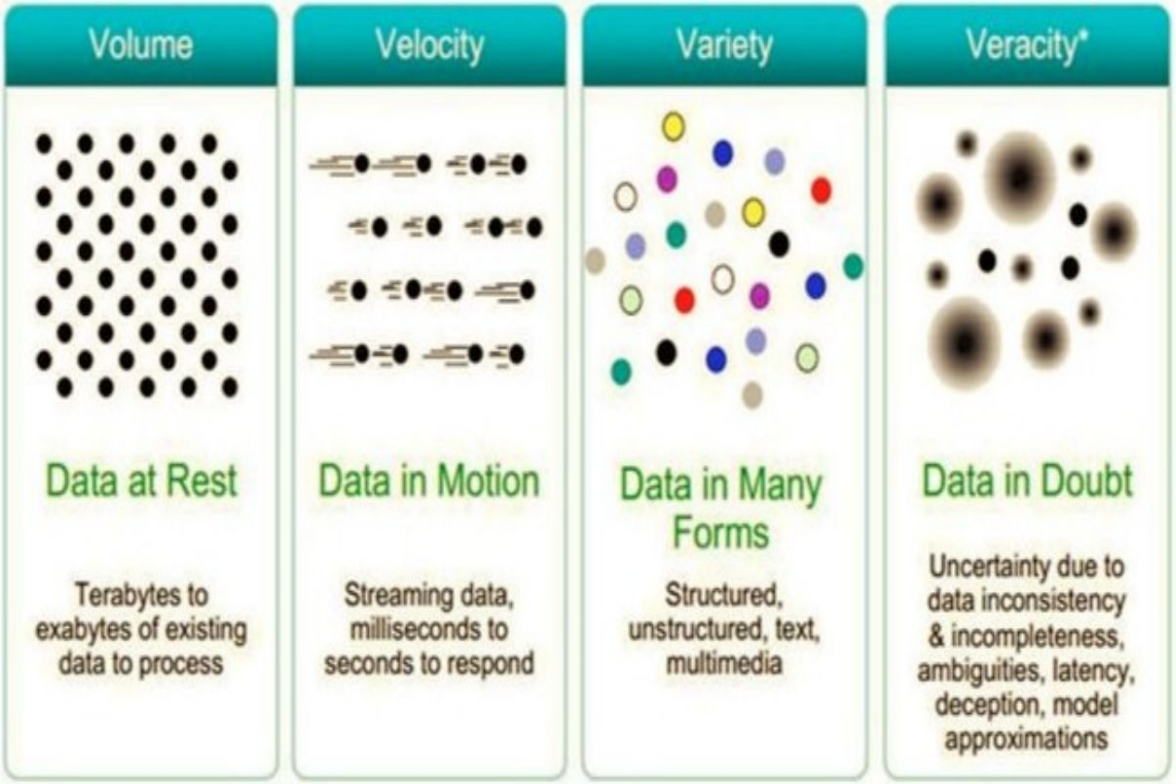

Nearly everyone across the organization engages with software

Yet, fewer than 25% of workers have access to analytical insights

BI Questions

Developing a business intelligence strategy is an important first step in implementing a BI solution.

- ✔ ***Who are the key stakeholders? Who will be using this system?***
- ✔ ***What departments need business intelligence and what will be measured?***
- ✔ ***What support do content authors and information consumers need?***
- ✔ ***Focus on Questions That Are Aligned With Your Business Strategy***
- ✔ ***Ask BI Questions That Give You Actionable Insights***
- ✔ ***Questions that Identify Opportunities for ROI***
- ✔ ***Questions That Identify Opportunities For New Sources of Revenue***
- ✔ ***Questions Identifying Cost-cutting Opportunities***
- ✔ ***Questions That Begin With "Why"***
- ✔ ***How Can You Understand Your Customers Better?***

As technology companies like Amazon, Meta, and Google continue to grow and integrate with our lives, they are leveraging big data technologies to monitor sales, improve supply chain efficiency and customer satisfaction, and predict future business outcomes.

Currently, there is so much big data that International Data Corporation (IDC) predicts the "Global Datasphere" will grow from 33 Zettabytes (ZB) in 2018 to 175 ZB in 2025. That's equal to a trillion gigabytes.

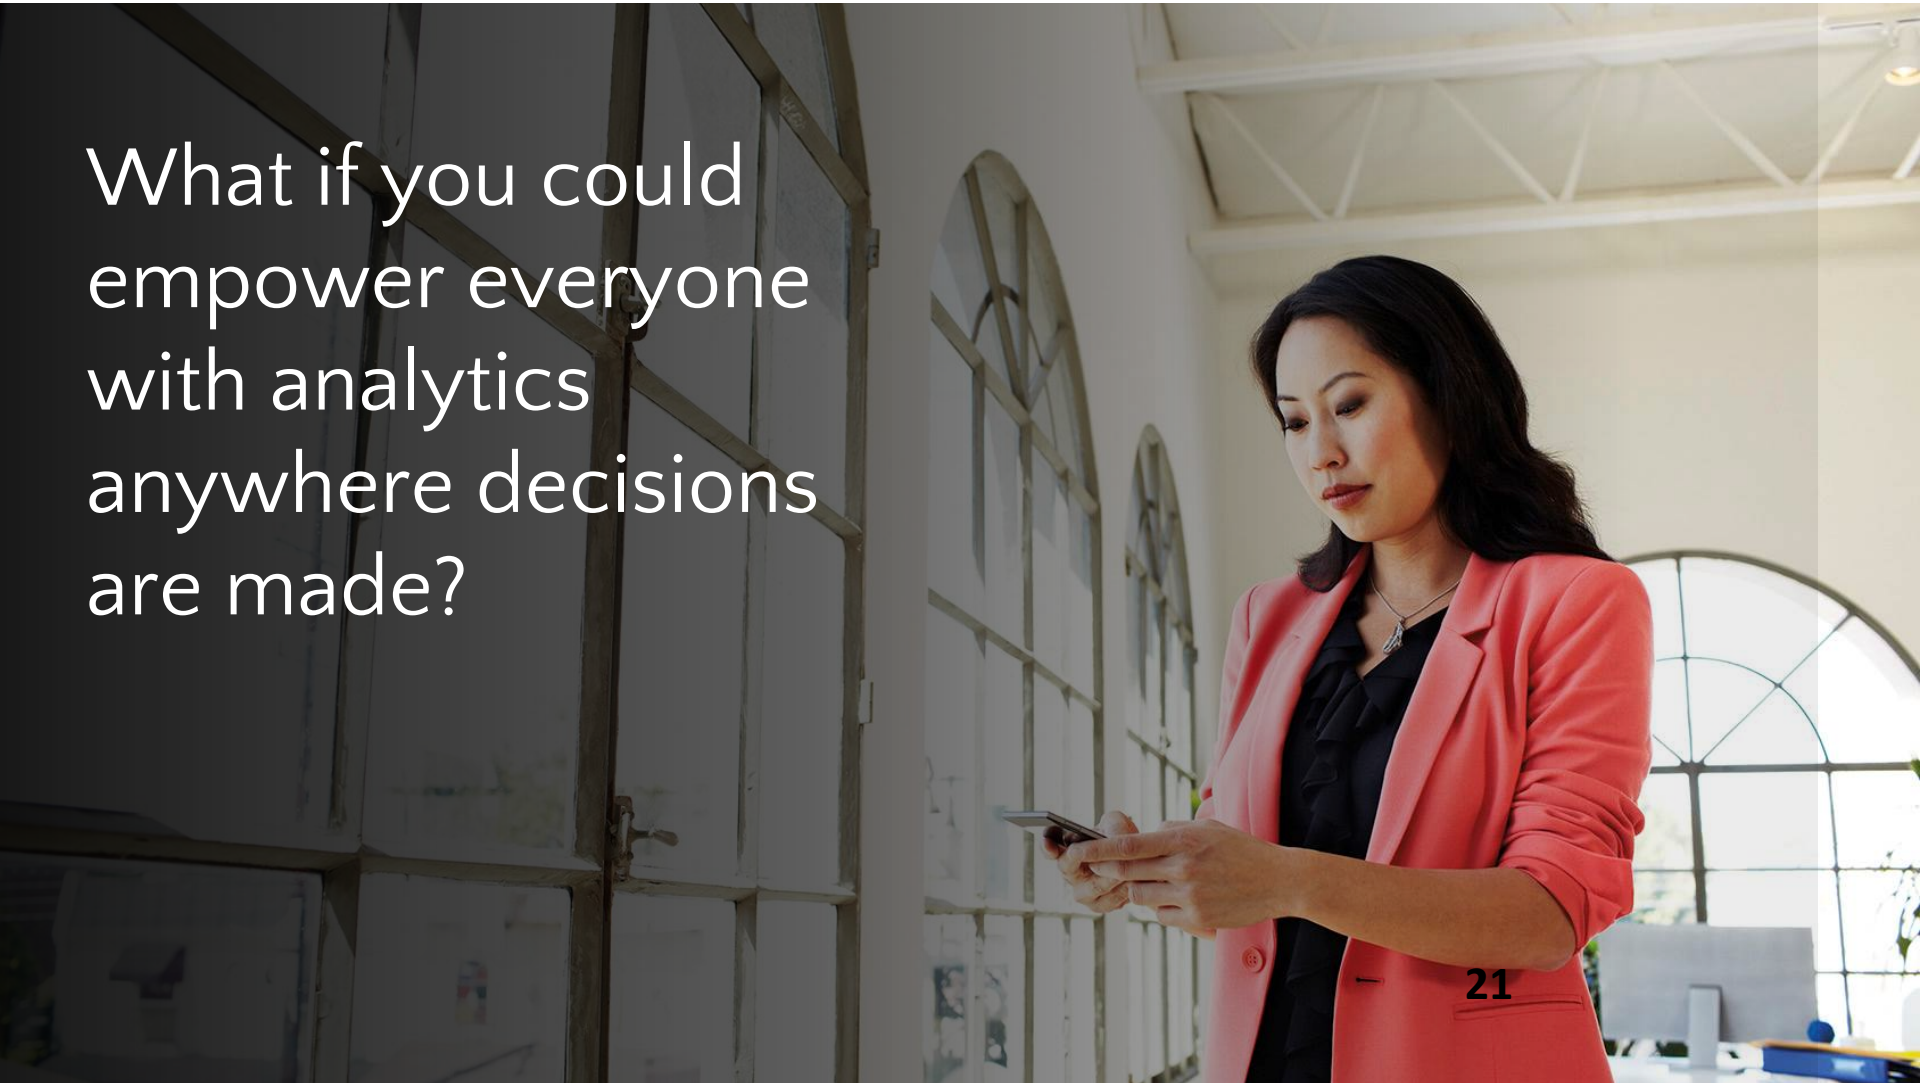
What if you could empower everyone with analytics anywhere decisions are made?

Today, BI extends to everyone

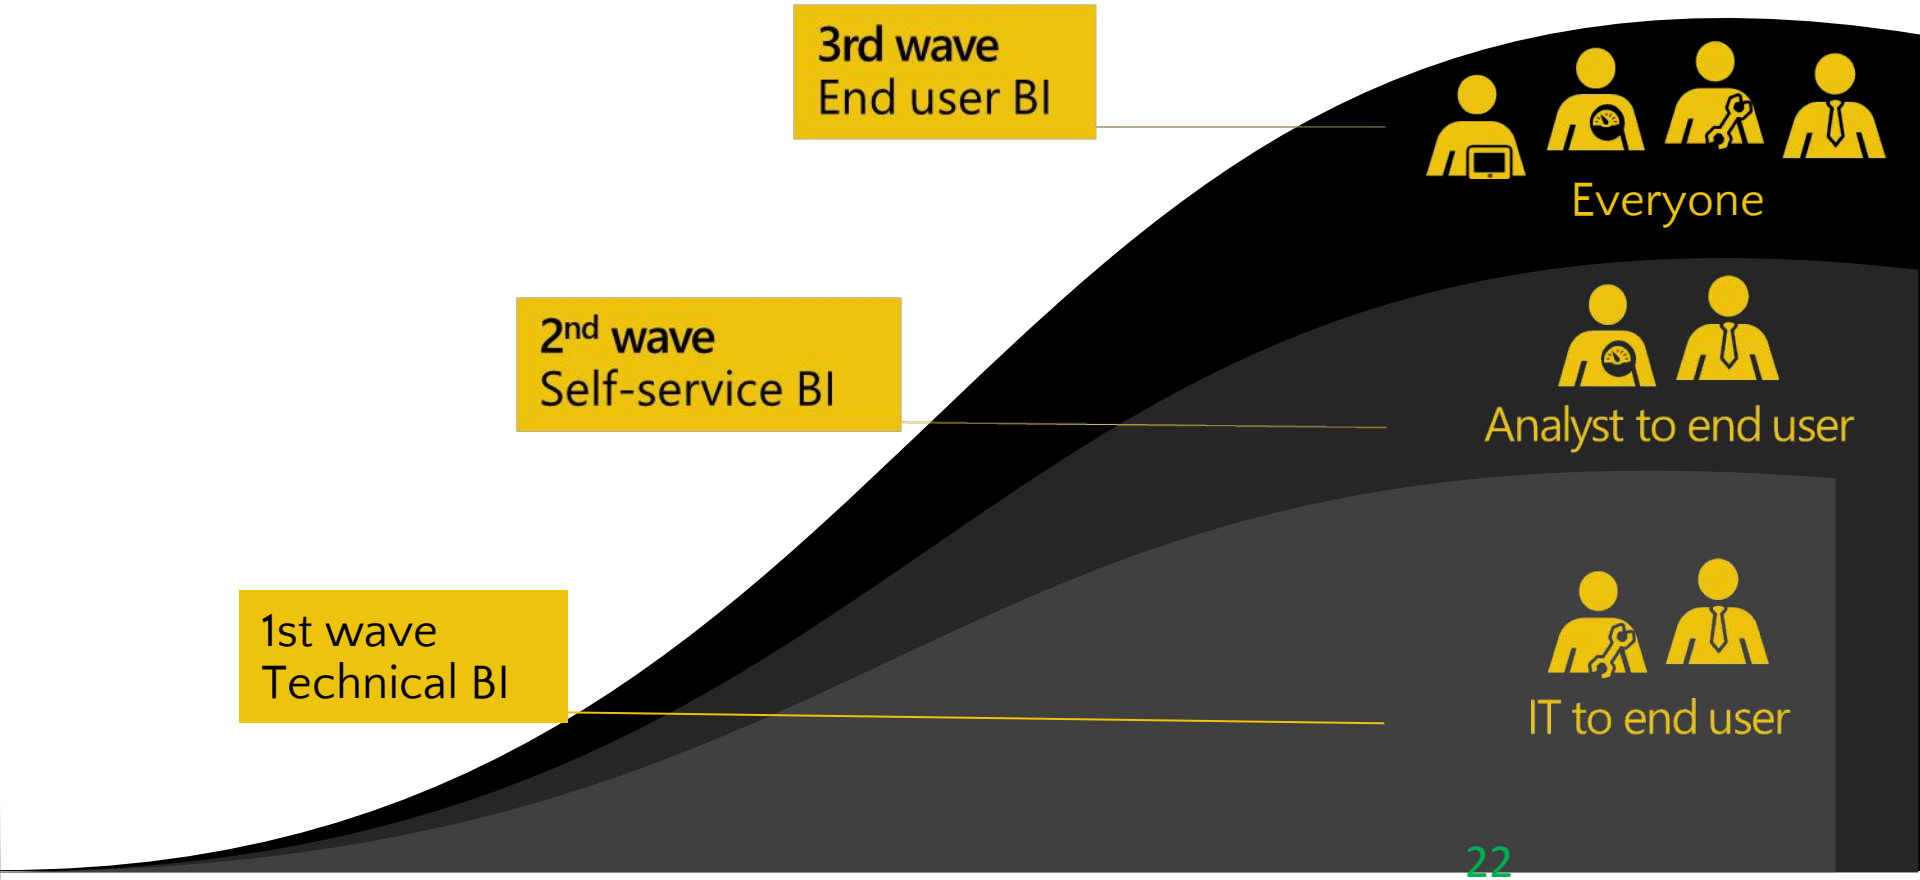

Two Data Factors

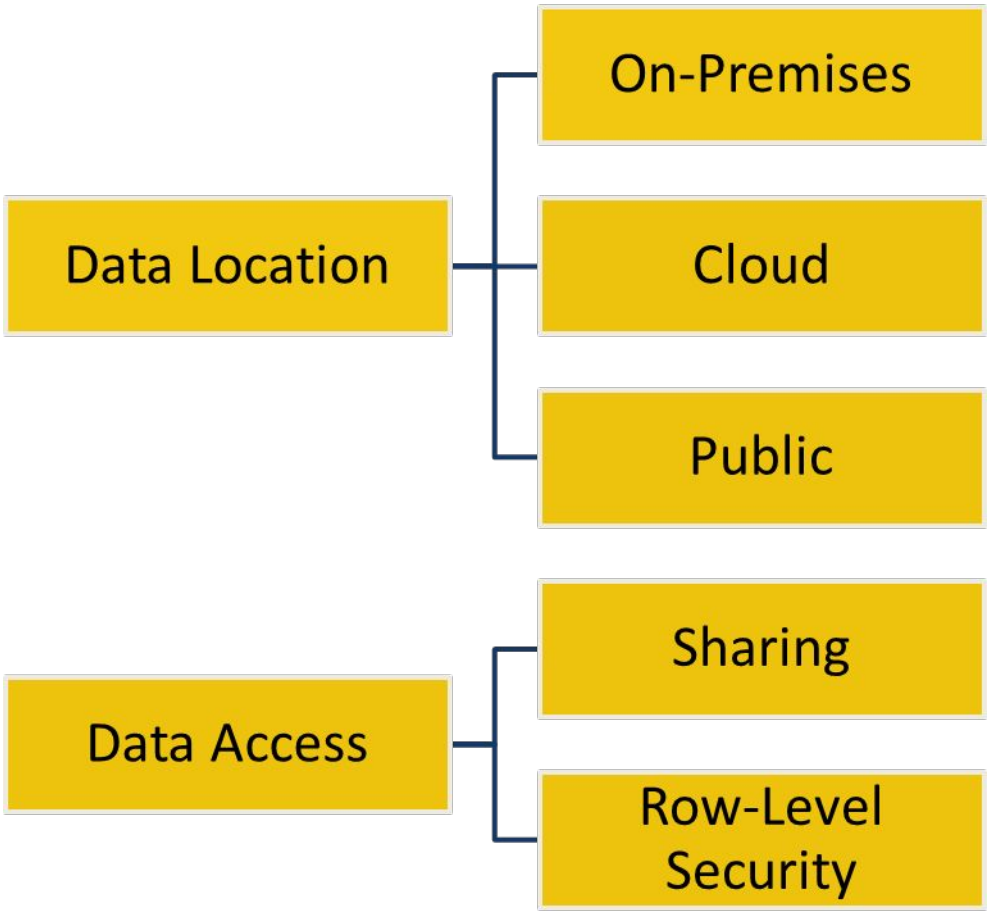

The Good

Experiments, observations, and numerical simulations in many areas of science and business are currently generating terabytes of data, and in some cases are on the verge of generating petabytes and beyond. Analyses of the information contained in these data sets have already led to major breakthroughs in fields ranging from genomics to astronomy and high-energy physics and to the development of new information-based industries.
- ***Frontiers in Massive Data Analysis, National Research Council of the National Academies***

The Bad

Given a large mass of data, we can by judicious selection construct perfectly plausible unassailable theories—all of which, some of which, or none of which may be right. - ***Paul Arnold Srere***

DATA ANALYTICS vs BUSINESS ANALYTICS

- Data analytics is a broad umbrella for finding insights in data
- Data analytics can refer to any form of analysis of data—whether in a spreadsheet, database, or app—where the intent is to uncover trends, identify anomalies, or measure performance.
- Additional mathematics or IT skills can help data analysts do everything from managing a database of subscribers to calculating yields for a potential investment.
- Data analytics (DA) is the technical process of mining data, cleaning data, transforming data, and building the systems to manage data. Data analytics takes large quantities of data to find trends and solve problems. Data analytics is not just confined to business applications—it's used across disciplines, from the government to science.

- Business analytics focuses on identifying operational insights.
- Business analytics focuses on the overall function and day-to-day operation of the business.
- A business analyst would deal less with the technical aspects of analysis and more with the practical applications of data insights.
- Some job responsibilities might include creating a streamlined workflow or choosing the best vendors.
- Business analytics (BA)refers to the process of taking your company's raw data and turning it into useful information, including identifying trends, predicting outcomes, and more.

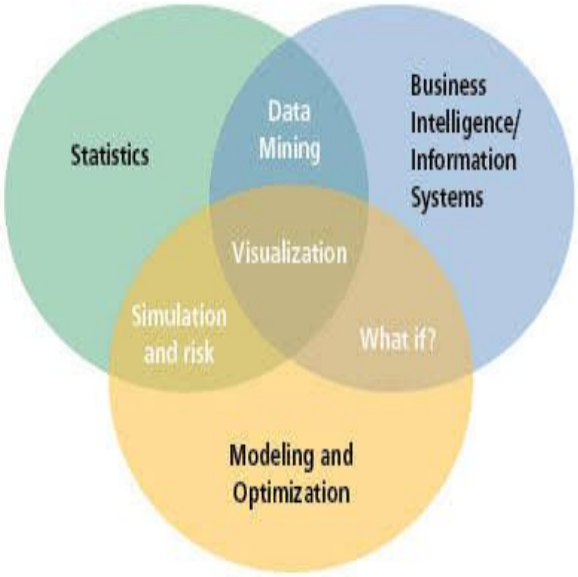

DIFFERENCE BETWEEN BUSINESS ANALYTICS AND DATA ANALYSIS

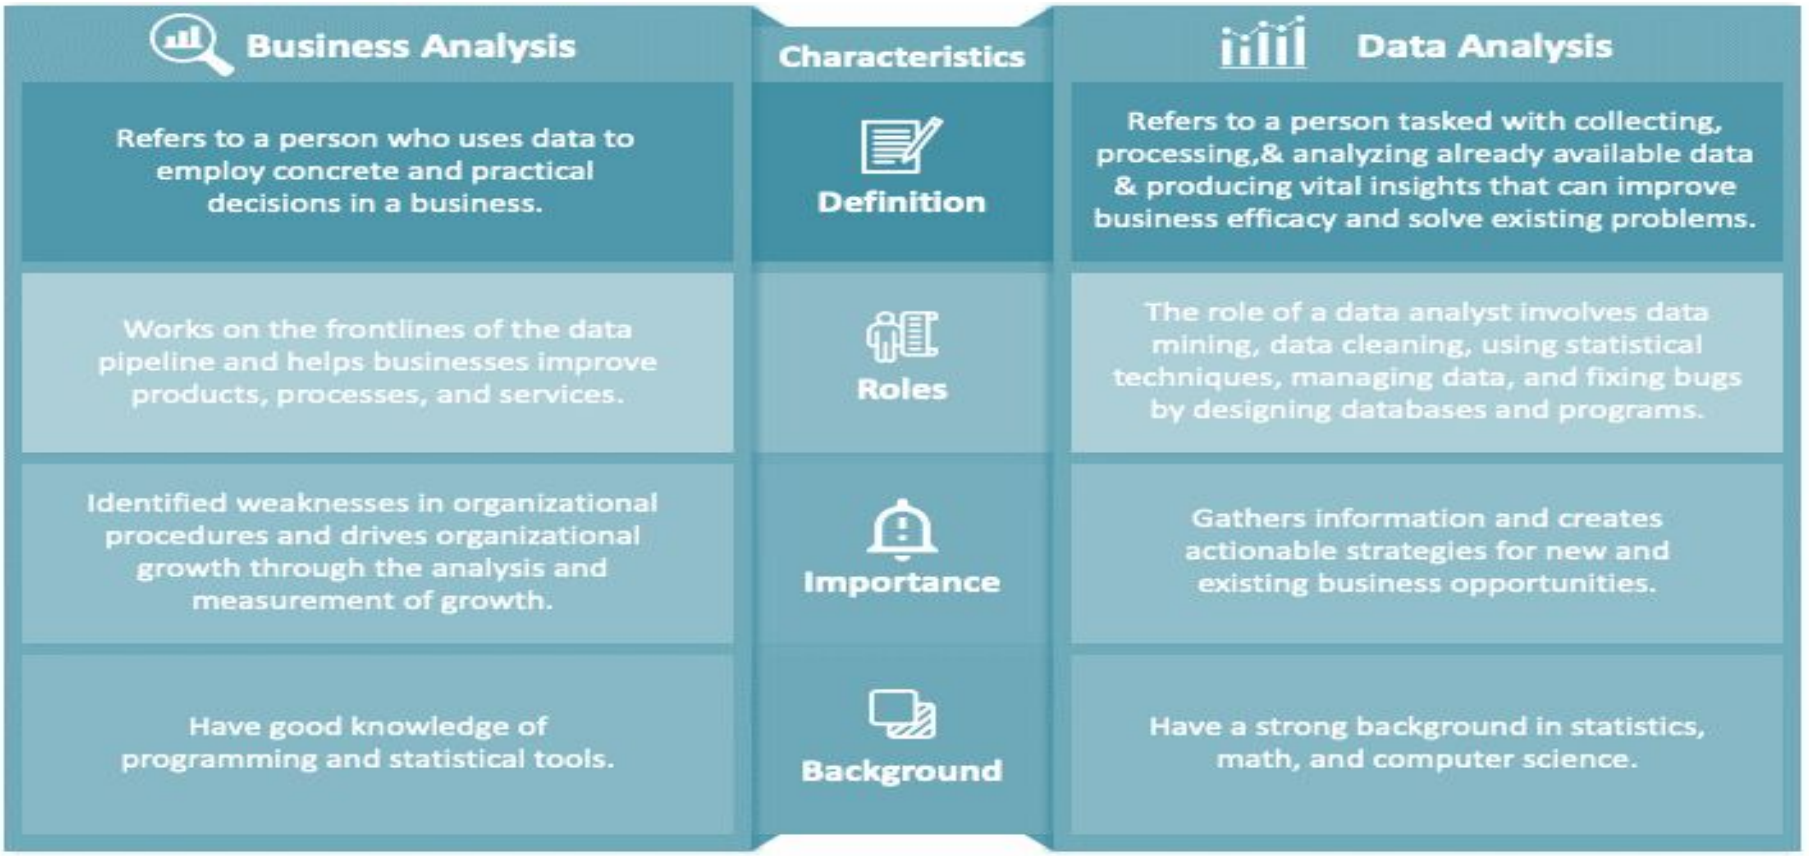

DIFFERENCE BETWEEN DATA ANALYTICS AND BIG DATA ANALYTICS

	DATA ANALYTICS	BIG DATA
NATURE	Like a book where you can find a solution to your problems.	considered as a Big Library where all the answers to all the questions are there but difficult to find the answers to your questions
STRUCTURE OF DATA	mostly structured data. It analyses the structured data to answer complex business queries, find solutions to business challenges, etc.	unstructured and raw data. The main aim of big data is to convert the raw data into meaningful data.
TOOLS	simple tools for statistical modelling and predictive modelling because the data to analyze is already structured and not complicated.	sophisticated technological tools such as automation tools or parallel computing tools to manage the Big Data
TYPE OF INDUSTRY	Used by IT Industries, Travel Industries, and Healthcare Industries. Data Analytics helps these industries to create new developments which are done by using historical data and analyzing past trends & patterns.	Used by industries such as banking industries, retail industries and many more. Big Data helps these industries in many ways to take some strategic business decisions.

Evolution of BI

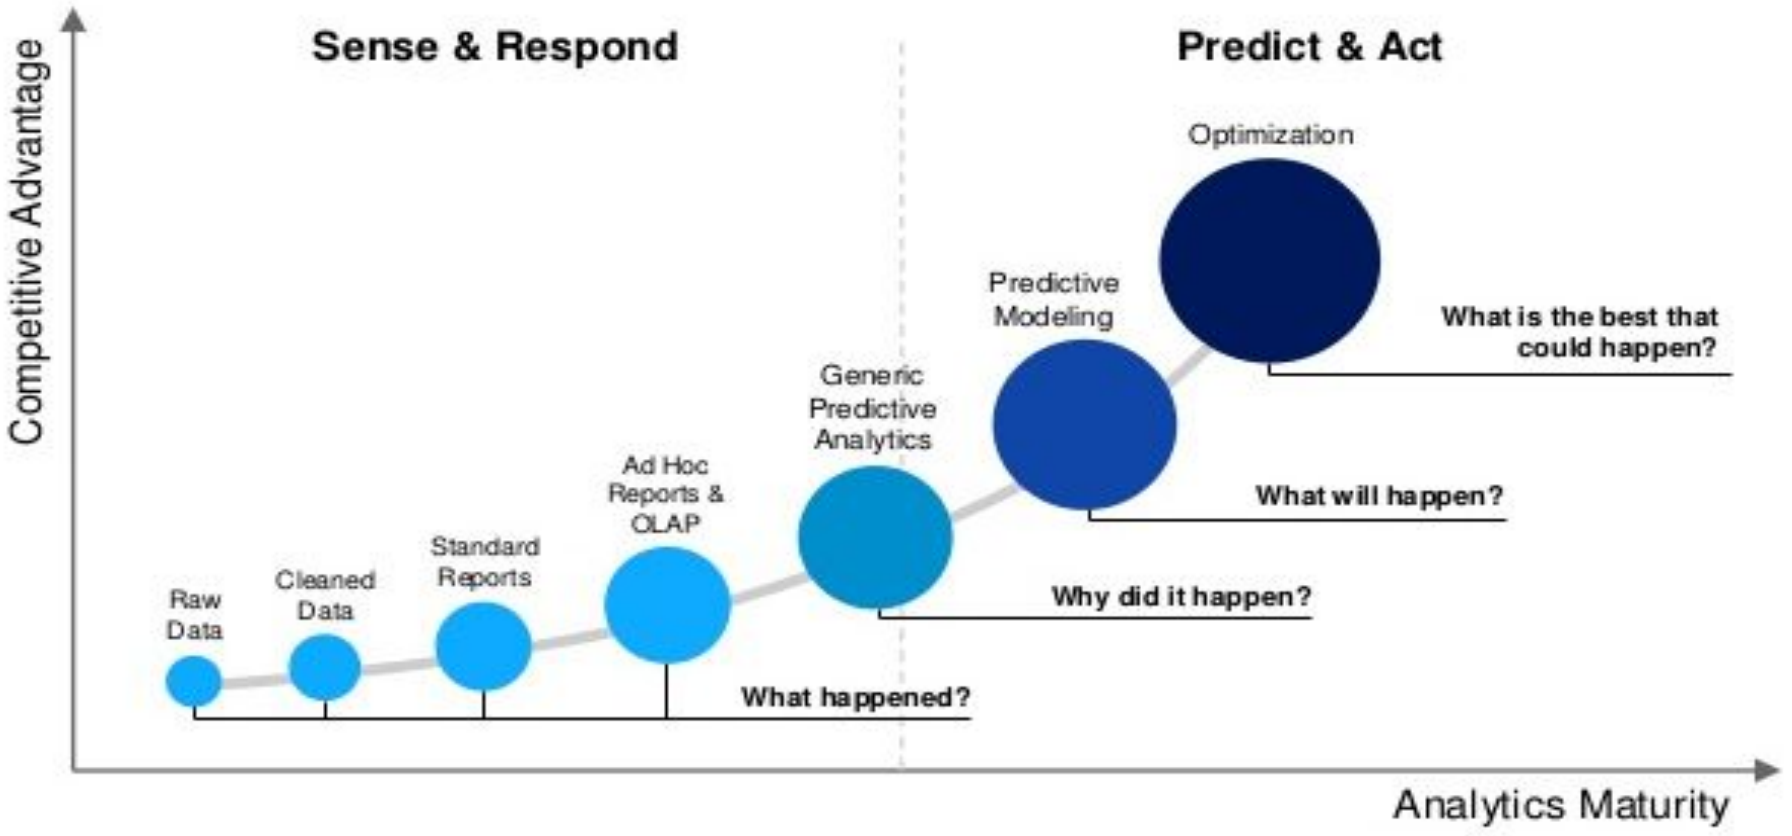

Characteristics of Data for Good Decision Making

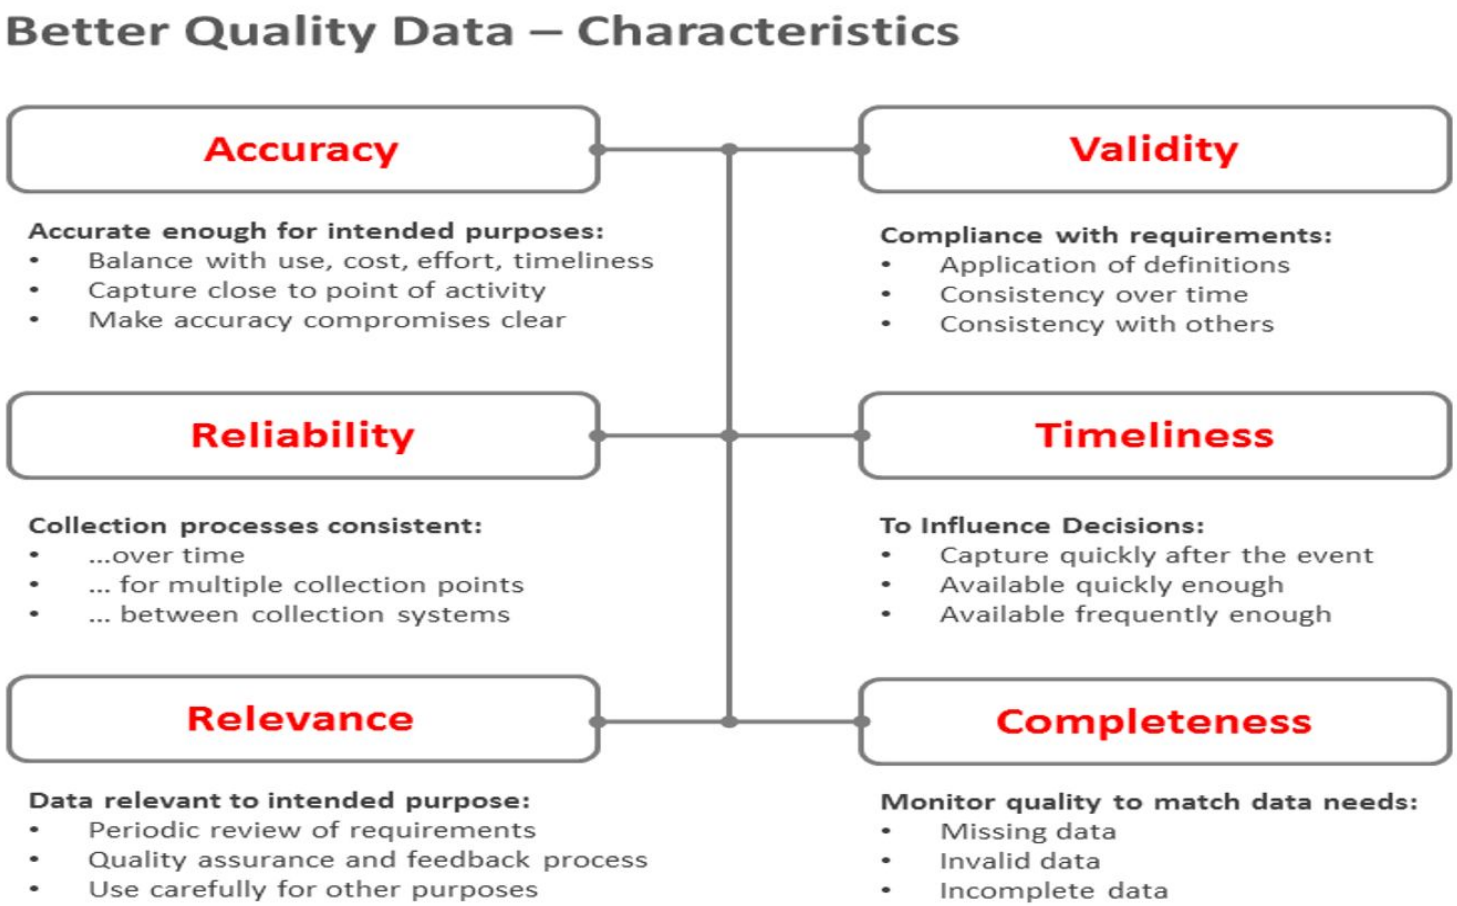

CRISP – DM

Cross-Industry Standard Process for Data Mining

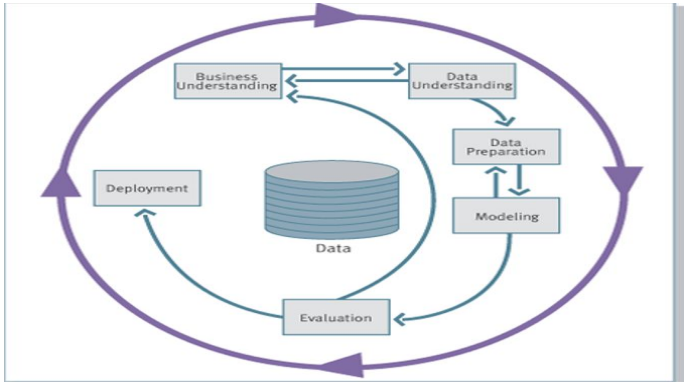

Business Understanding

Determine Business Objectives
Background
Business Objectives
Business Success Criteria

Situation Assessment
Inventory of Resources
Requirements, Assumptions, and Constraints
Risks and Contingencies
Terminology
Costs and Benefits

Determine Data Mining Goal
Data Mining Goals
Data Mining Success Criteria

Produce Project Plan
Project Plan
Initial Asessment of Tools and Techniques

Data Understanding

Collect Initial Data
Initial Data Collection Report

Describe Data
Data Description Report

Explore Data
Data Exploration Report

Verify Data Quality
Data Quality Report

Data Preparation

Data Set
Data Set Description

Select Data
Rationale for Inclusion / Exclusion

Clean Data
Data Cleaning Report

Construct Data
Derived Attributes
Generated Records

Integrate Data
Merged Data

Format Data
Reformatted Data

Modeling

Select Modeling Technique
Modeling Technique
Modeling Assumptions

Generate Test Design
Test Design

Build Model
Parameter Settings
Models
Model Description

Assess Model
Model Assessment
Revised Parameter Settings

Evaluation

Evaluate Results
Assessment of Data Mining Results w.r.t. Business Success Criteria
Approved Models

Review Process
Review of Process

Determine Next Steps
List of Possible Actions
Decision

Deployment

Plan Deployment
Deployment Plan

Plan Monitoring and Maintenance
Monitoring and Maintenance Plan

Produce Final Report
Final Report
Final Presentation

Review Project
Experience Documentation


Business Intelligence
Improves the decision making process
Makes it easy to Access and Share information
Helps you know your business
Enables real-time analysis with Quick Navigation
Reduces the risk of bottlenecks
Helps identify waste in the system


Descriptive
Describes what happened.
Diagnostic
Describes why something happened.
Predictive
Describes what's going to happen.
Prescriptive
What we should do about it.

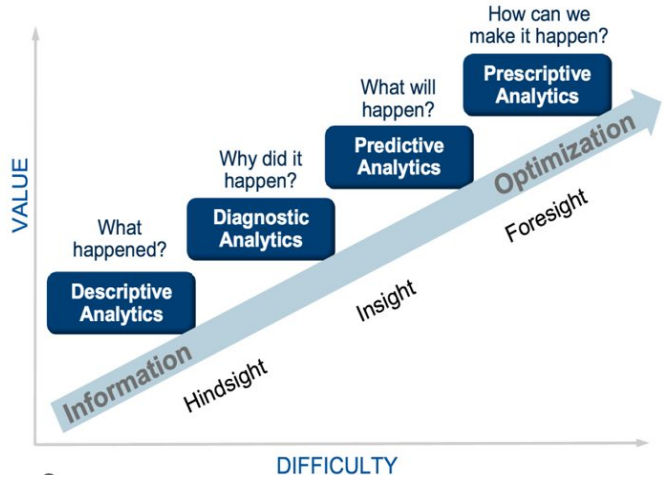
How can we make it happen?
Prescriptive Analytics
What will happen?
Predictive Analytics
Why did it happen?
Diagnostic Analytics
What happened?
Descriptive Analytics
VALUE
DIFFICULTY
Optimization
Foresight
Insight
Hindsight
Information

BIG DATA SYSTEMS PERSPECTIVE…

Systems Perspective: Processing Data

In-Memory Processing

Characteristics:

- Data processed directly in RAM, avoiding disk I/O.
- Extremely fast (low latency).

Advantages:

- Ideal for real-time analytics and low-latency applications (e.g., Spark, Apache Flink).
- Supports iterative algorithms and machine learning.

Challenges:

- Limited by available RAM.
- More expensive compared to disk-based solutions.

Systems Perspective: Processing from Secondary Storage

Characteristics:

- Data processed from hard drives or SSDs.
- Disk I/O introduces latency.

Advantages:

- Can handle massive datasets that don't fit in memory.
- Suitable for batch processing (e.g., Hadoop MapReduce).

Challenges:

- Slower compared to in-memory processing.
- Requires optimized data locality and access patterns.

Systems Perspective: Processing over the network

Characteristics:

- Distributed processing across multiple nodes in a network.
- Data often stored in HDFS, S3, or similar distributed systems.

Advantages:

- Scalability: Can process petabytes of data by leveraging many nodes.
- Redundancy: Fault-tolerant systems with data replication.

Challenges:

- Network latency can become a bottleneck.
- Requires efficient task scheduling and data shuffling (e.g., Apache Hadoop, Spark).

LOCALITY REFERENCES…

Locality of Reference: Principle & Examples

- - **Definition**: Locality of reference is a principle in computing that describes how programs tend to access a relatively small portion of their address space at any given time.
- - **Types of Locality**:
- - **Temporal Locality**: Recently accessed data is likely to be accessed again soon.
- - **Spatial Locality**: Data near a recently accessed location is likely to be accessed soon.
- - **Examples**:
- - **Code**: Sequential instruction execution (loops).
- - **Data**: Consecutive array accesses in loops.
- - **Memory Allocation**: Reuse of stack/heap data.

Impact of Locality on Performance

- - **Reduced Latency**:
- - Data in the CPU cache is faster to access than RAM or disk.
- - Better locality = fewer cache misses = reduced latency.
- - **Optimized Resource Usage**:
- - CPU pipelines stay efficient.
- - Reduced memory bandwidth contention.
- - **Examples**:
- - Loop unrolling and blocking in matrix multiplication.
- - Optimized database query plans.

Algorithms & Data Structures
Leveraging Locality

- - **Sorting Algorithms**:
- - Merge Sort benefits from spatial locality during merging phases.
- - **Search Structures**:
- - B-Trees/B+ Trees: Designed for efficient disk access.
- - **Dynamic Programming**:
- - Uses temporal locality by storing reusable subproblem solutions.

Data Organization for Better Locality

- - **On-Disk Data Layout**:
- - Contiguous allocation for files (e.g., ext4, NTFS).
- - Index structures like clustered B-Trees for databases.
- - **In-Memory Data Structures**:
- - Arrays vs. Linked Lists: Arrays have better spatial locality.
- - Cache-Aware Algorithms: Tailored for specific cache sizes.

Mitigating Latency Through Locality Optimization

- - **Software Optimizations**:
- - Code restructuring to improve data locality.
- - Using cache-friendly algorithms (e.g., blocking in matrix operations).
- - **Hardware Optimizations**:
- - Multi-level caches (L1, L2, L3).
- - Prefetching mechanisms.
- - **Real-World Applications**:
- - High-performance computing.
- - Database systems optimized for query efficiency.